Exercises for the course
Machine Learning 2
Summer semester 2020

Abteilung Maschinelles Lernen
Institut für Softwaretechnik und theoretische Informatik
Fakultät IV, Technische Universität Berlin
Prof. Dr. Klaus-Robert Müller
Email: klaus-robert.mueller@tu-berlin.de

Exercise Sheet 1

Exercise 1: Symmetries in LLE (25 P)

The Locally Linear Embedding (LLE) method takes as input a collection of data points $\vec{x}_1, \ldots, \vec{x}_N \in \mathbb{R}^d$ and embeds them in some low-dimensional space. LLE operates in two steps, with the first step consisting of minimizing the objective

$$\mathcal{E}(w) = \sum_{i=1}^{N} \left\| \vec{x}_i - \sum_j w_{ij} \vec{x}_j \right\|^2$$

where w is a collection of reconstruction weights subject to the constraint $\forall i : \sum_j w_{ij} = 1$, and where $\sum_j$ sums over the K nearest neighbors of the data point $\vec{x}_i$. The solution that minimizes the LLE objective can be shown to be invariant to various transformations of the data.

Show that invariance holds in particular for the following transformations:

(a) Replacement of all $\vec{x}_i$ with $\alpha \vec{x}_i$, for an $\alpha \in \mathbb{R}^+ \setminus \{0\}$,

(b) Replacement of all $\vec{x}_i$ with $\vec{x}_i + \vec{v}$, for a vector $\vec{v} \in \mathbb{R}^d$,

(c) Replacement of all $\vec{x}_i$ with $U\vec{x}_i$, where U is an orthogonal $d \times d$ matrix.

Exercise 2: Closed form for LLE (25 P)

In the following, we would like to show that the optimal weights w have an explicit analytic solution. For this, we first observe that the objective function can be decomposed as a sum of as many subobjectives as there are data points:

$$\mathcal{E}(w) = \sum_{i=1}^{N} \mathcal{E}_i(w) \qquad \text{with} \quad \mathcal{E}_i(w) = \left\| \vec{x}_i - \sum_j w_{ij} \vec{x}_j \right\|^2$$

Furthermore, because each subobjective depends on different parameters, they can be optimized independently. We consider one such subobjective and for simplicity of notation, we rewrite it as:

$$\mathcal{E}_i(w) = \left\| \vec{x} - \sum_{j=1}^{K} w_j \vec{\eta}_j \right\|^2$$

where $\vec{x}$ is the current data point (we have dropped the index i), where $\eta = (\vec{\eta}_1, \ldots, \vec{\eta}_K)$ is a matrix of size $K \times d$ containing the K nearest neighbors of $\vec{x}$, and w is the vector of size K containing the weights to optimize and subject to the constraint $\sum_{j=1}^{K} w_j = 1$.

(a) *Prove* that the optimal weights for $\vec{x}$ are found by solving the following optimization problem:

$$\min_w \quad w^\top C w \qquad \text{subject to} \quad w^\top \mathbb{1} = 1.$$

where $C = (\mathbb{1}\vec{x}^\top - \eta)(\mathbb{1}\vec{x}^\top - \eta)^\top$ is the covariance matrix associated to the data point $\vec{x}$ and $\mathbb{1}$ is a vector of ones of size K.

(b) *Show* using the method of Lagrange multipliers that the minimum of the optimization problem found in (a) is given analytically as:

$$w = \frac{C^{-1}\mathbb{1}}{\mathbb{1}^\top C^{-1}\mathbb{1}}.$$

(c) *Show* that the optimal w can be equivalently found by solving the equation $Cw = \mathbb{1}$ and then rescaling w such that $w^\top \mathbb{1} = 1$.

Exercise 3: t-SNE and Kullback-Leibler Divergence (25 P)

The t-SNE embedding algorithm operates by minimizing the Kullback-Leibler divergence between two discrete probability distributions p and q representing the input space and the embedding space respectively. These discrete distributions assign to each pair of data points (i, j) in the dataset the probability scores p_{ij} and q_{ij} respectively, corresponding to how close the two data points are in the input and embedding spaces. Once the exact probability functions are defined, the embedding algorithm proceeds by optimizing the function:

$$\begin{aligned} C &= D_{\mathrm{KL}}(p \,\|\, q) \\ &= \sum_{i=1}^{N} \sum_{j=1}^{N} p_{ij} \log\left(\frac{p_{ij}}{q_{ij}}\right) \end{aligned}$$

where p and q are subject to the constraints $\sum_{i=1}^{N} \sum_{j=1}^{N} p_{ij} = 1$ and $\sum_{i=1}^{N} \sum_{j=1}^{N} q_{ij} = 1$. Specifically, the algorithm minimizes q which itself is a function of the coordinates in the embedded space. Optimization is typically performed using gradient descent.

In this exercise, we derive the gradient of the Kullback-Leibler divergence, first with respect to the probability scores q_{ij}, and then with respect to the embedding coordinates of which q_{ij} is a function.

(a) *Show* that

$$\frac{\partial C}{\partial q_{ij}} = -\frac{p_{ij}}{q_{ij}}. \tag{1}$$

(b) The probability matrix q is now reparameterized using a 'softargmax' function:

$$q_{ij} = \frac{\exp(z_{ij})}{\sum_{k=1}^{N} \sum_{l=1}^{N} \exp(z_{kl})}$$

The new variables z_{ij} can be interpreted as unnormalized log-probabilities. *Show* that

$$\frac{\partial C}{\partial z_{ij}} = -p_{ij} + q_{ij}. \tag{2}$$

(c) *Explain* which of the two gradients, (1) or (2), is the most appropriate for practical use in a gradient descent algorithm. Motivate your choice, first in terms of the stability or boundedness of the gradient, and second in terms of the ability to maintain a valid probability distribution during training.

(d) The scores z_{ij} are now reparameterized as

$$z_{ij} = -\|\vec{y}_i - \vec{y}_j\|^2$$

where the coordinates $\vec{y}_i, \vec{y}_j \in \mathbb{R}^h$ of data points in embedded space now appear explicitly. *Show* using the chain rule for derivatives that

$$\frac{\partial C}{\partial \vec{y}_i} = \sum_{j=1}^{N} 4\,(p_{ij} - q_{ij}) \cdot (\vec{y}_i - \vec{y}_j).$$

Exercise 4: Programming (25 P)

Download the programming files on ISIS and follow the instructions.